Big Data Analytics Project

Notescape: A Big Data–Aware AI Learning Workspace

By

- **Mahnum Zahid (BSDSF22A022)**
- **Lubaba Khalid (BSDSF22A025)**
- **Hajra Amir (bsdsf22a027)**
- **Rabbiya Bukhari (bsdsf22044)**

1. Introduction

1.1 Background

With the growth of digital education and online learning platforms, there is an increasing need for tools that can efficiently handle large volumes of unstructured data, such as PDFs, scanned notes, and lecture slides. Traditional learning systems are not well-equipped to handle this influx of information, particularly when it comes to continuous interaction and AI-driven personalization.

NoteScape is designed as an AI-powered learning workspace that integrates Big Data principles to manage, process, and analyze educational content. By applying scalable storage, distributed processing, and intelligent retrieval systems, NoteScape aims to enhance the learning experience and provide personalized recommendations based on user interaction

1.2 Problem Statement

Current educational tools are not optimized for dealing with large datasets. They face several limitations:

- Inefficient handling of large documents: Traditional systems struggle to process and manage large volumes of unstructured content.

- Limited support for semantic search: Most systems rely on keyword-based searches, which are inadequate for extracting meaningful insights from large, diverse datasets.

- Lack of long-term learning analytics: There is no real-time analysis of user behavior and no mechanism for providing long-term insights into students' learning progress.

- Poor scalability for concurrent users: Traditional learning systems are often unable to scale effectively to support large numbers of users, leading to performance bottlenecks.

These limitations highlight the need for a Big Data–aware system that can support distributed data processing, offer scalable storage solutions, and enable intelligent content retrieval.

1.3 Project Objectives

The main objectives of the NoteScape project are:

- AI-powered personal learning workspace: Develop a learning environment that leverages AI for personalized learning support.

- Large-scale document processing: Implement systems to efficiently handle large volumes of educational content, such as PDFs and scanned documents.

- Retrieval-Augmented Generation (RAG): Apply RAG to enable intelligent document retrieval and contextual generation of responses based on study materials.

- Learning behavior tracking: Collect and analyze user interactions over time to improve learning outcomes.

- Demonstration of Big Data principles: Apply concepts such as distributed storage, parallel processing, and scalability to a real-world educational platform.

2. Project Overview

NoteScape is an AI-driven platform designed to enhance the learning experience by leveraging Big Data principles. The system is capable of processing large study materials, enabling semantic search, generating AI-powered flashcards, and providing analytics to monitor and improve user learning behaviors.

The design of the system ensures that it is scalable and fault-tolerant, essential for handling high volumes of data and multiple concurrent users. The platform's main features include intelligent content retrieval, AI-driven flashcard generation, and personalized learning analytics, making it a robust solution for modern educational environments.

3. Core Features

The core features of NoteScape are integral to its functionality as an AI-powered, Big Data-aware learning system:

- AI-powered flashcard generation: Using machine learning models, NoteScape can automatically generate flashcards from study materials to aid in revision and long-term memory retention.

- Retrieval-Augmented Generation (RAG): RAG allows the system to retrieve relevant sections of study material based on a user's query, enhancing the accuracy and context of AI-generated responses.

- Semantic search: NoteScape includes a powerful semantic search engine that allows users to search for content not just by keywords but also by meaning, improving search relevance.

- Asynchronous document processing: The system is designed to handle document processing (such as OCR and chunking) in the background without interrupting the user experience, ensuring smooth and efficient performance.

- Learning analytics dashboard: This feature tracks and visualizes user progress, helping both students and educators identify areas for improvement and optimize learning strategies.

- Secure user authentication: Firebase authentication ensures that user data is protected and provides secure access to the platform.

These features require integration with Big Data technologies such as distributed storage for documents, vector databases for embeddings, and asynchronous processing for non-blocking tasks, making NoteScape a highly scalable and efficient platform.

4. Functional and Non-Functional Requirements

4.1 Functional Requirements

The functional requirements for the system are designed to ensure a smooth and effective learning experience:

- Document upload and processing: The system must be capable of handling large PDFs, scanned notes, and lecture slides, converting them into a structured format for analysis.
- Flashcard generation: Automatically generate flashcards based on the content uploaded by the user, enabling spaced repetition for enhanced learning.
- Spaced repetition tracking: Support for the "Again / Hard / Good / Easy" spaced repetition system to optimize study efficiency.
- Learning behavior tracking: Monitor user interactions and study patterns to provide personalized learning insights and track weaknesses over time.

4.2 Non-Functional Requirements

To support Big Data workloads, the system ensures:

- Scalability: The platform should be able to handle a large number of concurrent users, particularly as the volume of data grows.

- Low latency: The system must provide real-time responses for user interactions, ensuring minimal delays.
- Fault tolerance: The system should be resilient to failures, with background tasks that are capable of recovering from errors without affecting the user experience.
- Data consistency: Ensures that all data across the system is consistent and accurate, particularly for user analytics.
- Security: User data, including study materials and personal information, should be securely stored and protected using industry-standard encryption and access controls.

5. Need for Big Data Technologies

The scale of NoteScape's operations necessitates the use of Big Data technologies:

- Large unstructured documents: PDFs, scanned images, and other unstructured formats require specialized tools for efficient processing and storage.
- High-dimensional embeddings: Embeddings generated from the text data (e.g., using machine learning models) require a storage solution capable of handling high-dimensional vector data.
- Continuous study interaction events: User behavior, such as study time, flashcard reviews, and learning patterns, must be tracked in real-time.
- Long-term learning analytics: NoteScape aims to provide long-term learning insights, which requires a system capable of analyzing vast quantities of historical data.

- Multiple concurrent users: The system must handle multiple users simultaneously without performance degradation, requiring distributed systems and parallel processing.

Traditional relational databases are not equipped to manage these challenges effectively, making Big Data technologies essential.

6. Big Data Characteristics (4Vs)

Volume

NoteScape generates large datasets that include study notes, document chunks, flashcards, and real-time analytics events. These datasets must be processed efficiently using scalable systems.

Variety

The data in NoteScape comes in various formats, including:

- Text: Raw study material (PDFs, lecture slides)
- Vectors: Generated embeddings for semantic search and RAG
- Metadata: Information about users, their study habits, and performance
- Analytical events: Data points representing user interactions with the platform (e.g., time spent studying, flashcard reviews)

Velocity

NoteScape handles real-time interactions such as flashcard reviews and learning progress updates, requiring a system that can handle high-speed data flow and provide immediate feedback to users.

Veracity

The accuracy of AI-generated responses and learning analytics is critical. The system must ensure that the data processed is reliable, and that AI outputs are consistent and valid, particularly when providing learning recommendations.

7. Data Storage Strategy (Polyglot Persistence)

Polyglot Persistence allows the use of multiple storage technologies, each suited to specific types of data. NoteScape uses the following components:

- Firebase Auth: For secure user authentication and identity management.
- PostgreSQL: For core data storage, including user profiles, metadata, and flashcards.
- pgvector: A PostgreSQL extension used for storing document embeddings, which are essential for RAG and semantic search.
- Amazon S3: For storing large documents, such as PDFs and scanned notes.
- Redis: Used optionally for caching frequently accessed data and enhancing system performance.
- Background Workers: Handle asynchronous tasks such as OCR, chunking, and embedding generation, ensuring that long-running tasks do not block user interactions.

This strategy ensures that each data type is stored in the most efficient manner possible, improving the overall performance and scalability of the system.

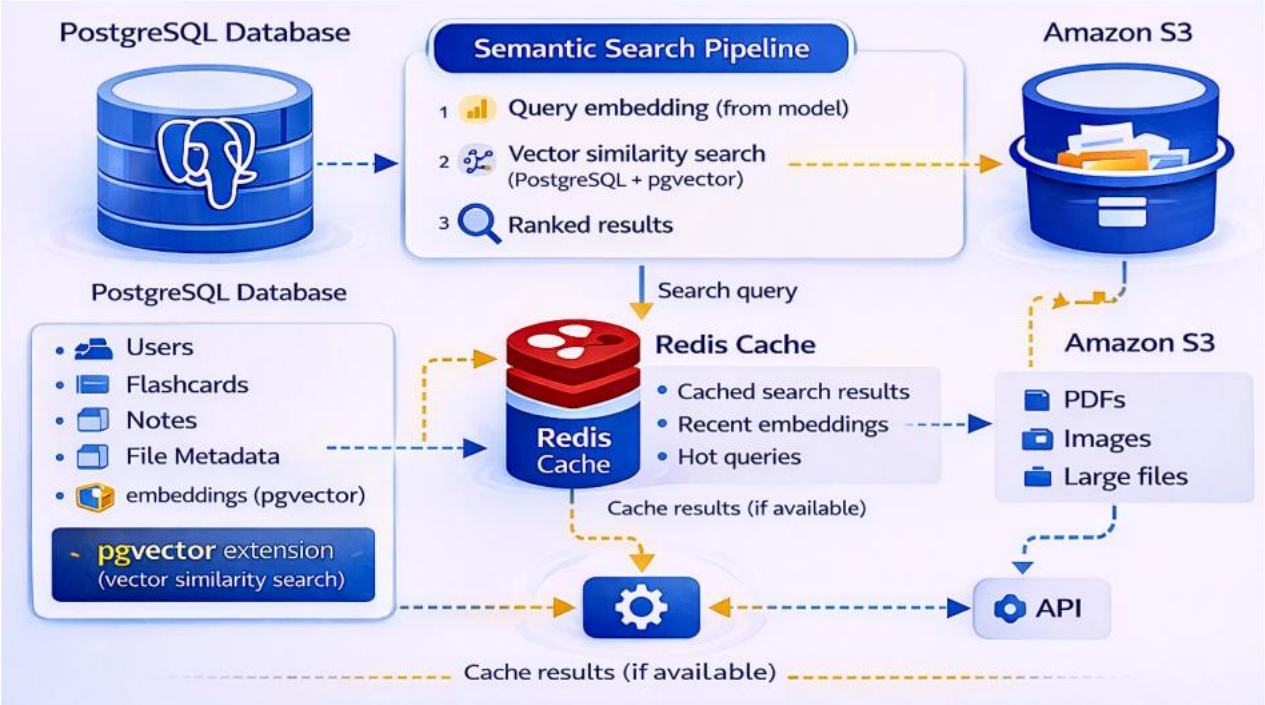

8. Asynchronous and Distributed Processing

To ensure efficient processing of large datasets, the system employs asynchronous and distributed processing. Long-running tasks such as OCR, document chunking, and embedding generation are handled by background workers. The benefits of this approach include:

- Non-blocking user experience: Users can interact with the system while background tasks are running, ensuring a seamless experience.
- Parallel task execution: Tasks like document processing and embedding generation are executed in parallel, improving performance and reducing latency.
- Fault isolation: Background tasks are isolated from the main application, ensuring that errors in one process do not disrupt the entire system.

9. Document Chunking Strategy

In order to efficiently process large documents, NoteScape employs a chunking strategy. Documents are split into smaller, more manageable chunks, which are processed independently. This approach allows:

- Parallel embedding generation: By processing smaller chunks in parallel, the system can generate embeddings faster.
- Reduced memory overhead: Smaller chunks use less memory, allowing for more efficient processing.
- Improved retrieval accuracy: Smaller, more specific chunks can be retrieved based on their relevance to the user's query.

Chunking is a fundamental strategy in Big Data processing as it enables the system to handle large volumes of data while minimizing resource usage.

10. Retrieval-Augmented Generation (RAG)

RAG allows NoteScape to generate intelligent, context-aware responses to user queries. The process involves:

1. Chunking and embedding documents: Documents are split into smaller chunks and converted into embeddings.
2. Retrieving relevant chunks using vector similarity: The system retrieves the most relevant chunks based on semantic similarity to the user’s query.
3. Passing retrieved context to the language model: The context from the relevant chunks is passed to a language model, such as GPT, to generate an accurate response.

4. Generating context-aware responses: The response generated by the model is grounded in the retrieved data, ensuring it is contextually relevant.

This approach ensures that the AI-generated content is always accurate and relevant, making it an essential feature for personalized learning.

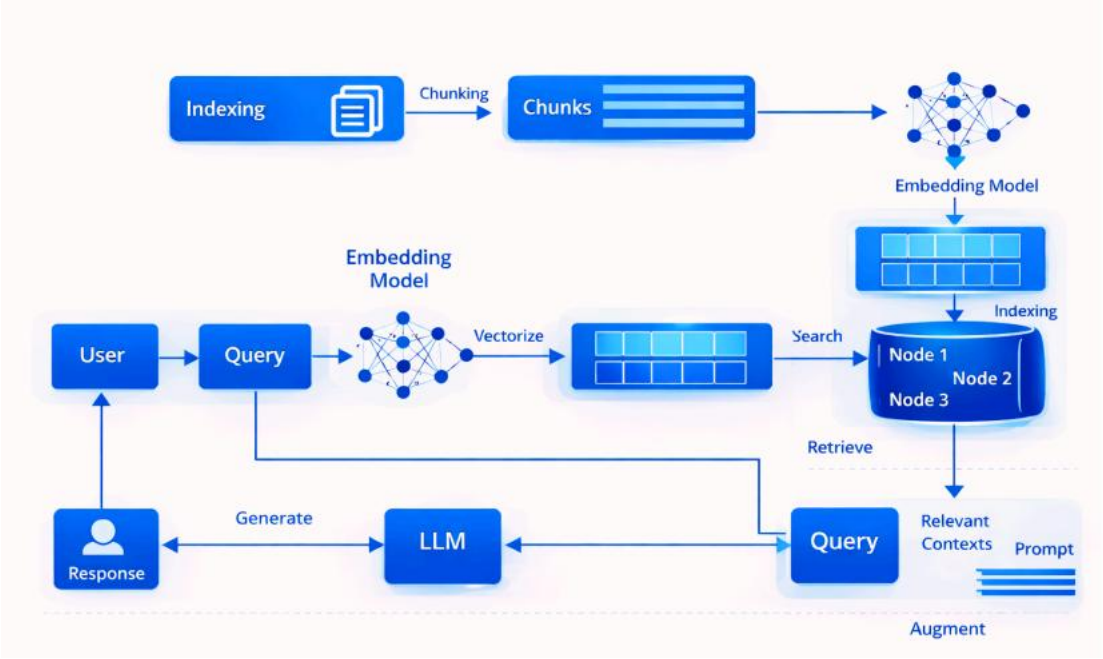

11. RAG Processing Pipeline

The RAG pipeline is integral to the operation of NoteScape and includes the following steps:

1. Document upload: Users submit their study materials, which are processed by the system.

2. OCR & text extraction: Text is extracted from scanned documents.

3. Chunking: The document is split into smaller, manageable chunks.

4. Embedding generation: Each chunk is converted into a vector embedding.

5. Vector retrieval: The system retrieves the most relevant chunks based on similarity to the user's query.

6. AI-generated responses: The system generates a response or flashcard based on the retrieved data.

Each of these steps demonstrates the Big Data principles applied to enhance the learning experience.

12. System Architecture Overview

The overall architecture of NoteScape consists of several key components:

- Frontend (User Interface): Provides the user interface for interacting with the system.

- Firebase Authentication: Manages user authentication and identity securely.

- Backend API: Handles the core functionality, such as interacting with databases and performing AI tasks.

- Asynchronous Workers: Responsible for background tasks like OCR, chunking, and embedding generation.

- PostgreSQL with pgvector: Stores core data (e.g., user profiles, study data) and vector embeddings.

- Amazon S3: Used for storing large documents, such as PDFs and scans.

Each component contributes to the system's scalability, fault tolerance, and performance, ensuring an optimal experience for users.

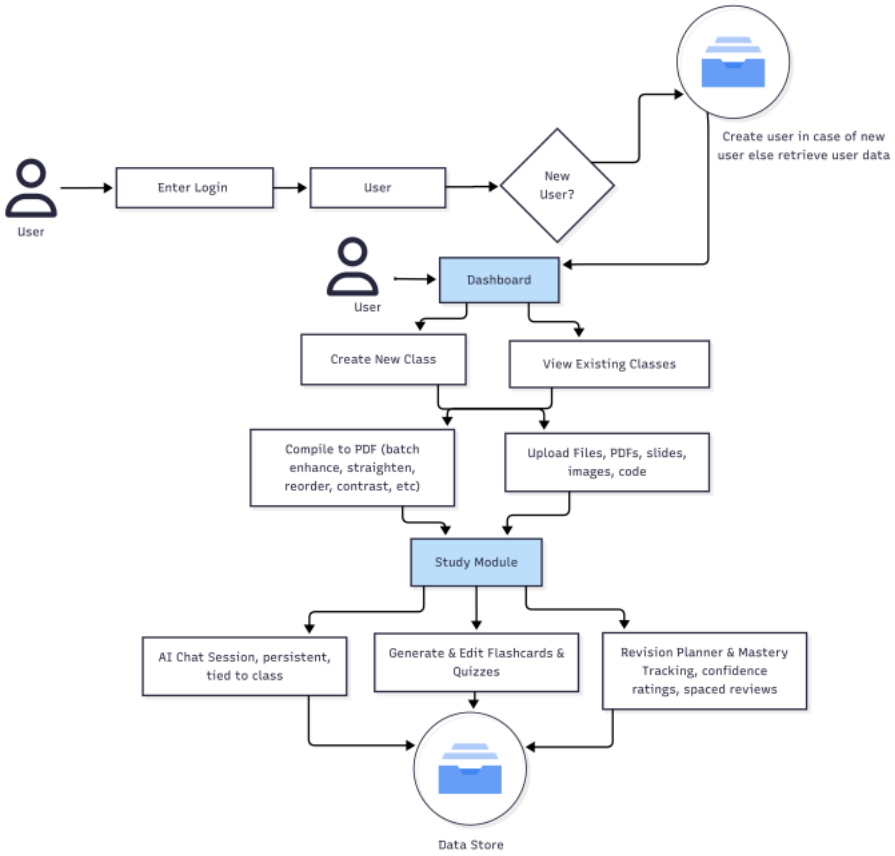

13. Conclusion and Future Work

Conclusion

The NoteScape project successfully applies Big Data principles to build an AI-powered learning platform. It integrates various technologies such as polyglot persistence, asynchronous processing, and RAG-based AI to provide a scalable, efficient, and intelligent learning experience.

Future Enhancements

Future enhancements include:

- Qdrant integration: For large-scale vector search and improved retrieval capabilities.
- Concept graph visualization: To enhance the representation and understanding of learning materials.

- Advanced learning recommendations: Powered by AI, to offer personalized learning paths and resources.

NoteScape has set the foundation for future Big Data applications in educational technology.